

# Homework 5: Image Classification

December 1, 2023

**Due Date:** December 15 by 23:59:59

## Introduction

In this assignment, you will implement and test various image classification models on the [CIFAR-10](#) dataset. The goals of this assignment are as follows:

- Implement and compare the linear classifier and the full-connected neural network.
- Implement and compare the SVM loss and the cross-entropy loss.
- Implement and compare the AdamW and the SGD optimizer.
- Implement and compare the StepLR and the CosineAnnealingLR scheduler.
- Train and test two types of classifiers.

You can learn how to create, train, and test a model using PyTorch [here](#). You are highly encouraged to go through [this tutorial](#) before you start.

Here are some other supplementary materials that may help you:

- [PyTorch Documentation](#)
- [PyTorch Chinese Documentation](#)
- [Dive into deep learning](#)

## 1 Define Classifiers (20 pts.)

### 1.1 Linear classifier (10 pts.)

Add your own code to the **LinearClassifier** class to define a linear classifier. Your classifier is required to process a mini-batch data.

### 1.2 Full-connected neural network classifier (10 pts.)

Add your own code to the **FCNN** class to define a full-connected neural network classifier. You are responsible for choosing the network depth, width, and activation type.

## 2 Define loss function (20 pts.)

You need to implement the SVM loss and the cross-entropy loss from scratch. The weight decay term for regularization is also required, but you are not required to implement it explicitly since you can use `weight_decay` defined in optimizer for this purpose.

### 2.1 SVM loss function (10 pts.)

Add your own code to the `svmloss()` function to define an SVM loss. You need to implement it from scratch instead of calling pre-implemented PyTorch functions that directly finish the task.

### 2.2 Cross-entropy loss function (10 pts.)

Add your own code to `crossentropyloss()` function to define cross-entropy loss. You also need to implement it from scratch instead of calling pre-implemented PyTorch functions that directly finish the task, such as `torch.nn.functional.cross_entropy` and `torch.nn.functional.nll_loss`.

## 3 Implement the training and testing function (30 pts.)

There is a whole training code in [PyTorch Tutorial: train a classifier](#), you can learn from it. In this task, you need to implement the `train()` and `test()` function that can choose a model, optimizer, scheduler, and so on; see the end of the `main.py` for details.

## 4 Compare AdamW and SGD optimizer (10 pts.)

Train the classifiers you implemented using the AdamW (`torch.optim.AdamW`) and SGD (`torch.optim.SGD`) optimizer and compare the loss and accuracy curves. Put the results in your report.

## 5 Compare SVM and Cross-entropy loss (10 pts.)

Train the classifiers you implemented using the SVM and Cross-entropy loss and compare the loss and accuracy. Put the loss and accuracy curves and the final classification accuracies in your report. You can use **TensorBoard** in PyTorch to record and visualize the loss and accuracy curves. Here is a [tutorial](#) introducing **TensorBoard**.

## 6 Compare StepLR and CosineAnnealingLR scheduler (10 pts.)

Train the classifiers you implemented using two learning rate schedulers, including the StepLR (`torch.optim.lr_scheduler.StepLR`) and CosineAnnealingLR (`torch.optim.lr_scheduler.CosineAnnealingLR`) scheduler and compare the loss and accuracy curves. Put the results in your report.

## 7 Report

You have now completed the entire process of this project. Put all the visualizations and results in your report. More experiments and discussions are encouraged.

## 8 Submit

Be sure to zip your code and final report; Name it as **StudentID\_YourName\_HW5.zip**.